

Database Systems - Project 2

Jjjj222

1 Overview

In this TinySQL DBMS program, all required features are fully supported, including insertion, deletion, selection, projection, sorting, duplication removal, and join, with unlimited size of tuples in any relations. All of the operations, especially selection and join, are optimized to reduce the overall disk I/Os. Moreover, this program also provides decent error detections: every invalid query will be stopped immediately, with the location of potential problems being pointed out clearly.

In addition, all features are designed to be as similar to MySQL as possible, such as table display, commands, and interactive user interface with autocomplete and history support. This program is also flexible enough to further support other queries which is not yet defined in current TinySQL grammar.

This report is organized as follow: the section 2 will provide the information of setup and usage; section 3 will explicate the structure of this program, and then the details of functionality and performance of every operation will be discussed in section 4, and the optimization details will be addressed in section 5.

2 README

This section provides instructions to build and execute TinySQL, as well as folder structure and development environment.

2.1 Build TinySQL

Type “make” under the project folder “tinysql/” to build the project. The executable file “tinysql” will be in “bin/” folder after compilation.

```
shell> cd tinysql
shell> make
```

2.2 Execute TinySQL

Interactive Mode: To execute TinySQL in interactive mode, type “./bin/tinysql” in folder “./tinysql”. The prompt “tinysql>” indicates that you can start using TinySQL now. If you want to stop TinySQL, use command “quit” to exit.

```
shell> ./bin/tinysql
tinysql>
...
tinysql> quit
Bye
```

Batch Mode: To execute TinySQL in batch mode, type “./bin/tinysql” and follow by the name of the file, in which are commands you want to execute. TinySQL will exit automatically after encountering errors or finishing execution.

```
shell> cd tinysql
shell> ./bin/tinysql $FILE
```

All the results will be display on screen by standard output. If you want to dump the results to a file, please redirect the output.

```
shell> ./bin/tinysql $FILE > $OUT_FILE
```

If you want the results to be both on screen and in a file, please use linux pipe and tee

```
shell> ./bin/tinysql $FILE | tee $OUT_FILE
```

To run multiple commands without leaving, use the “source” command in interactive mode.

```
shell> ./bin/tinysql
tinysql> source $FILE
...
tinysql>
```

2.3 Project Structure

The project structure is as follows:

- makefile
- bin/: executable
- db_sim_api/: StorageManager library
- dep/: dependency files generated by gcc
- include/: some additional header files
- obj/: object files generated by gcc
- results/: results of automatic test by “make test”.
- src/: source codes
- testcases/: some testcases. Note that those with suffix *.in will be executed by typing “make test”, and the results will be compared with *.out.

2.4 Development Environment

This program was developed under both Linux and MacOS X in c++11 and c99.

- OS: Ubuntu 12.04.5 LTS
- compiler: gcc 4.8.1
- lex: 2.5.35
- yacc: GNU Bison 2.5
- make: GNU Make 3.81
- libraries: The GNU Readline Library

3 Program Architecture

The followings are basic architecture of the TinySQL program:

3.1 Command Line Interface

TinySQL supports two kinds of user interface mode: interactive mode and batch mode. I use GNU readline library to support history and autocomplete in the interactive mode. In batch mode, either by specifying dofile or by “source” command, queries will be executed consecutively, and the prompt will become the file name and line number of the current command. The commands other than TinySQL syntax, such as “source” and “quit”, are implemented here.

file: cmd.cpp, cmd.h

3.2 TinySQL Parser

I used lex/yacc to parse the TinySQL commands and build the parsing tree. The structure of the parsing tree “**struct** tree_node” is as follows:

```
struct tree_node
{
    int type;
    const char* value;
    struct tree_node* next;
    struct tree_node* child;
};
```

file: parser.c, parser.h, sql.l, sql.y

3.3 Query Processor

The main functionalities of TinySQL are implemented in here. All the queries which has been transformed to parsing tree will be processed by routines in “**class** QueryMgr”. Logic query plans will be built for the “SELECT” queries: it’s also a tree-like structure, with additional information, and was implemented in “**class** QueryNode”. After being constructed, the logic query plan will then be optimized and transformed to another logic query plan with estimated fewer overall disk I/Os. The physical query plan will reuse the same tree structure, carrying out the query by choosing desirable functions.

Another tree structure “**class** ConditionNode” will be built under “WHERE” conditions and used for filtering out tuples. After construction, the “**class** ConditionMgr” can take a tuple as the parameter, extract corresponding data from it, and then execute the “**class** ConditionNode”. The result “**true**” mean this tuple fit the “WHERE” condition; “**false**”, otherwise.

file: query.cpp, query.h

3.4 Database Interface

Most of the utilities for accessing StorageManager library are in “**class** `hwMgr`”. This class is charge of initiating `MainMemory`, `Disk` and `SchemaManager`, as well as managing and allocating their resources.

Other classes such as “**class** `TinyRelation`” and “**class** `TinyTuple`” are wrappers that combine low level operations into high level operations. They work with “**class** `RelScanner`”, “**class** `RelWriter`”, and **class** `RelSorter` to read/write or sort relations.

class RelScanner: This class is the main routine to scan tables. It will take a relation and the currently available memory spaces for initialization; the the range to scan can be further specified. After initialization, users can use `RelScanner::get_next()` function to iterate through the table. The `RelScanner` will keep iterators to the current positions of both disk and memory, and automatically load blocks from disk to memory while it running out of tuples in memory.

class RelSorter: This class is the main routine for sorting relations. It also takes a relation and the currently available memory spaces for initialization, and the order by which tuples to be sorted. and then it will sorted the relation automatically by applying n-pass multiway merge, while the pass number depends on the size of the relation. Note that `RelSorter` will create temporary relations to store the partial sorted tuples.

class RelWriter: This class will load the last block in a relation to the memory, append the new tuples to it, and then write it back to the disk.

file: `dbMgr.cpp`, `dbMgr.h`, `wrapper.cpp`, `wrapper.h`

3.5 Utilities

These files contain some general utilities, such as those for error messege reporting, table drawing, string tokenizing, and debugging.

file: `obj_util.cpp`, `obj_util.h`, `tiny_util.cpp`, `tiny_util.h`, `debug.h`, `util.h`

4 Command and Query

This section will discuss the implementation detail of queries, along with their performance evaluations and examples.

4.1 CREATE TABLE

This command defines the schema of a relation. It will report error if a relation with the same name already exists in the database. There is no disk I/O required.

Example: Create a table “example” with 2 attributes, “attr0” and “attr1”.

```
tinysql> CREATE TABLE example (attr0 INT, attr1 STR20)
Query OK (0 disk I/O, 0 ms)

tinysql> show tables
+-----+
| Tables_in_tinysql |
+-----+
| example            |
+-----+
1 row in set (0 disk I/O, 0 ms)
```

4.2 DROP TABLE

This command will delete an existing table. It will report errors if the specified table does not exist in the database. There is no disk I/O required.

Example: Delete the table “example”.

```
tinysql> show tables
+-----+
| Tables_in_tinysql |
+-----+
| example            |
| example2           |
+-----+
2 rows in set (0 disk I/O, 0 ms)

tinysql> DROP TABLE example
Query OK (0 disk I/O, 0 ms)

tinysql> show tables
+-----+
| Tables_in_tinysql |
+-----+
| example2           |
+-----+
1 row in set (0 disk I/O, 0 ms)
```

4.3 INSERT INTO

The INSERT INTO command adds a new tuple into an existing relation. The number of disk I/O varies based on the current size of the relation: if the last block of the relation is not full, then TinySQL will load it to the memory first, and then write it back after appending the new tuple. It will cost total 2 disk I/Os. On the other hand, if the last block is empty, then TinySQL will overwrite it directly, with only 1 disk I/O.

In current implementation, TinySQL always appends the new tuple to the back of the relation, even if there might be some “hole” where the original tuple was deleted. Note that TinySQL doesn’t support NULL, due to the limitation of StoreManager library (It uses int instead of int* to store integer value). Every NULL value, including those

caused by attributes not being specified in the query , will be set to default value (0 for INT and empty string for STR20).

Example: Insert tuples into table “example”.

```
tinysql> CREATE TABLE example (attr0 INT, attr1 STR20)
Query OK (0 disk I/O, 0 ms)

tinysql> INSERT INTO example (attr0, attr1) VALUES (0, "A")
Query OK (1 disk I/O, 74.63 ms)

tinysql> INSERT INTO example (attr1, attr0) VALUES ("B", 1)
Query OK (2 disk I/O, 149.89 ms)

tinysql> SELECT * FROM example
+-----+-----+
| attr0 | attr1 |
+-----+-----+
|      0 | A     |
|      1 | B     |
+-----+-----+
2 rows in set (1 disk I/O, 75.52 ms)
```

4.4 DELETE FROM

This command delete tuples from relation. It will load, from disk to memory, all non-empty blocks in the relation, and then delete every tuple which matches the *WHERE* condition. If any tuple is deleted in an block, then the block will be written back to disk, otherwise the block in memory will be simply discarded. The total number of disk I/Os for a relation R would be $B(R)$ + number of blocks modified.

Example: Delete tuples from table “example”.

```
tinysql> SELECT * FROM example
+-----+-----+
| attr0 | attr1 |
+-----+-----+
|      0 | A     |
|      1 | B     |
+-----+-----+
2 rows in set (1 disk I/O, 75.52 ms)

tinysql> DELETE FROM example WHERE attr1="A"
Query OK (2 disk I/O, 149.78 ms)

tinysql> SELECT * FROM example
+-----+-----+
| attr0 | attr1 |
+-----+-----+
|      1 | B     |
+-----+-----+
1 row in set (1 disk I/O, 75.41 ms)
```

4.5 SELECT FROM

The SELECT command prints out tuples which comply with specified requirements. Before generating any outputs, TinySQL will first construct a logic query plan and then optimize it. A physical query plan will reuse the same tree structure, and calculate desired results with different algorithms. The functionalities and performances of different operations would be discussed in the following sections.

Example: A query with its corresponding logic query plan

```
SELECT example.attr0, example.attr1, example2.attr0
FROM example, example2
WHERE example.attr0 < example2.attr0
ORDER BY example.attr0

root:
  '- PROJECTION: {example.attr0, example.attr1, example2.attr0}
    '- ORDER_BY: example.attr0
      '- WHERE
        |- WHERE_OPTION
        |   '- COMP_OP: <
        |     |- COLUMN_NAME: example.attr0
        |     '- COLUMN_NAME: example2.attr0
        '- CROSS_PRODUCT
          |- BASE_NODE: example
          '- BASE_NODE: example2
```

Example2: A query with 200 tuples

```
tinysql> SELECT * FROM example200
...
200 rows in set (200 disk I/O, 14927 ms)
```

4.5.1 Projection (SELECT)

This operation simply scans the table and selects specified columns. It takes $B(R)$ disk I/Os to scan the table; However, if the results need to be materialized, it will take additional $B(R)$ disk I/Os to write the results into the temporary table..

Example: Project column “attr0”.

```
tinysql> SELECT * FROM example
+-----+-----+
| attr0 | attr1 |
+-----+-----+
|      0 | A      |
|      1 | B      |
+-----+-----+
2 rows in set (1 disk I/O, 75.15 ms)

tinysql> SELECT attr0 FROM example
+-----+
| attr0 |
+-----+
```

	0	
	1	
+-----+		
2 rows in set (1 disk I/O, 75.52 ms)		

4.5.2 Duplication Removal (DISTINCT)

This operation removes duplicated tuples. TinySQL uses sort-based algorithm: first sort the relation if it has not yet been sorted, and then scan the relation and output only the tuple which is not the same as the next tuple. Therefore, it will take $B(R)$ disk I/Os with a potential sorting to complete its job.

Example: Remove duplication tuples

tinysql> SELECT * FROM example		
+-----+		
attr0	attr1	
+-----+		
0	A	
0	A	
1	B	
+-----+		
3 rows in set (1 disk I/O, 74.78 ms)		
tinysql> SELECT DISTINCT * FROM example		
+-----+		
attr0	attr1	
+-----+		
0	A	
1	B	
+-----+		
2 rows in set (3 disk I/O, 224.67 ms)		

Example2: A query with 200 tuples

tinysql> SELECT DISTINCT * FROM example200		
...		
200 rows in set (1400 disk I/O, 104482 ms)		

4.5.3 Sorting (ORDER BY)

This operation uses multiway merge sort. It could be 1-pass, 2-pass, or n-pass, which is based on the size of the relation. If it takes more than 1 pass to complete the sort, TinySQL will create temporary relations to store the partially sorted tuples. It requires $2 \times B(R) \times n$ (pass) disk I/Os to sort a relation.

table size		1	2	3	4	5	6	7	8	9	10	20	30
disk I/O		3	6	9	12	15	18	21	24	27	50	100	150
table size		40	50	60	70	80	90	100	110	120			
disk I/O		200	250	300	350	400	630	700	770	840			

Table 1: disk I/Os with table size for 1 sorting + 1 scan

Example: Sorting

```
tinysql> SELECT * FROM example
+-----+-----+
| attr0 | attr1 |
+-----+-----+
|      0 | A     |
|      2 | C     |
|      1 | B     |
+-----+-----+
3 rows in set (1 disk I/O, 74.78 ms)

tinysql> SELECT * FROM example ORDER BY attr0
+-----+-----+
| attr0 | attr1 |
+-----+-----+
|      0 | A     |
|      1 | B     |
|      2 | C     |
+-----+-----+
3 rows in set (3 disk I/O, 224.67 ms)
```

Example2: A query with 200 tuples

```
tinysql> SELECT * FROM example200 ORDER BY attr0
...
200 rows in set (1400 disk I/O, 104483 ms)
```

4.5.4 Selection (WHERE)

The selection operation only filters out the output of other operations, it doesn't involve disk I/O.

Example: Selection

```
tinysql> SELECT * FROM example
+-----+-----+
| attr0 | attr1 |
+-----+-----+
|      0 | A     |
|      1 | B     |
+-----+-----+
2 rows in set (1 disk I/O, 75.52 ms)

tinysql> SELECT * FROM example WHERE attr1="A"
+-----+-----+
| attr0 | attr1 |
+-----+-----+
|      0 | A     |
+-----+-----+
1 row in set (1 disk I/O, 75.15 ms)
```

Example2: A query with 200 tuples

```
tinysql> SELECT * FROM example200 WHERE attr0=0
...
Empty set (200 disk I/O, 14926 ms)
```

4.5.5 Cross Product (FROM)

The cross product operation combine every tuple in the relations. It takes $B(R) \times B(S) / M$ disk I/Os, where M = available memory size. The physical query plan will assign 1 memory block to the larger relation, and ther rest memory blocks to the smaller relation.

Example: Cross product

```
tinysql> SELECT * FROM example
+-----+-----+
| attr0 | attr1 |
+-----+-----+
|      0 | A      |
|      1 | B      |
+-----+-----+
2 rows in set (1 disk I/O, 75.41 ms)

tinysql> SELECT * FROM example1
+-----+-----+
| attr0 | attr1 |
+-----+-----+
|      1 | A      |
|      0 | B      |
+-----+-----+
2 rows in set (1 disk I/O, 75.04 ms)

tinysql> SELECT * FROM example, example1
+-----+-----+-----+-----+
| example.attr0 | example.attr1 | example1.attr0 | example1.attr1 |
+-----+-----+-----+-----+
|              0 | A              |              1 | A              |
|              0 | A              |              0 | B              |
|              1 | B              |              1 | A              |
|              1 | B              |              0 | B              |
+-----+-----+-----+-----+
4 rows in set (3 disk I/O, 223.93 ms)
```

4.5.6 Natural Join (Theta Join)

The natural join operation only exists after optimization. A cross product with equation will be transformed to a natural join. Natural join will, based on the attributes in equation, sort relations first. And then it will scan both table and do cross product for only tuples with matched attribute. The disk I/O depends heavy on the values in the relations. However, most of the time, the time complexity will reduce from $O(B(R) \times B(S))$ of original cross product to $O(B(R) + B(S))$.

Example: Natural join (Theta join)

```

tinysql> SELECT * FROM example
+-----+-----+
| attr0 | attr1 |
+-----+-----+
|      0 | A      |
|      1 | B      |
+-----+-----+
2 rows in set (1 disk I/O, 75.41 ms)

tinysql> SELECT * FROM example1
+-----+-----+
| attr0 | attr1 |
+-----+-----+
|      1 | A      |
|      0 | B      |
+-----+-----+
2 rows in set (1 disk I/O, 75.04 ms)

tinysql> SELECT * FROM example, example1 WHERE example.attr0 = example1.
      attr0
+-----+-----+-----+-----+
| example.attr0 | example.attr1 | example1.attr0 | example1.attr1 |
+-----+-----+-----+-----+
|              0 | A              |              0 | B              |
|              1 | B              |              1 | A              |
+-----+-----+-----+-----+
2 rows in set (10 disk I/O, 746.34 ms)

```

5 Optimizaion

To speed up queries, TinySQL implements several optimizations.

5.1 Insert

TinySQL will calculate the location for the new tuple before accessing the disk: if the block to load from disk is empty, then TinySQL will overwrite it directly, save 1 disk I/O from loading it to memory..

Example: Overwrite directly

```

tinysql> SELECT * FROM example
+-----+-----+-----+-----+
| attr0 | attr1 | attr2 | attr3 |
+-----+-----+-----+-----+
|      0 |      0 |      0 |      0 |
+-----+-----+-----+-----+
1 row in set (1 disk I/O, 75.26 ms)

tinysql> INSERT INTO example (attr0, attr1, attr2, attr3) VALUES (0, 0, 0,
      0)
Query OK (2 disk I/O, 149.52 ms)

tinysql> INSERT INTO example (attr0, attr1, attr2, attr3) VALUES (1, 1, 1,
      1)

```

```
Query OK (1 disk I/O, 75.15 ms)
```

5.2 Scan

TinySQL keep track of the “holes” caused by deletion. If all the tuples in a block are deleted, then scanner will skip it.

Example: Skip empty block

```
tinysql> SELECT * FROM example
+-----+-----+-----+-----+
| attr0 | attr1 | attr2 | attr3 |
+-----+-----+-----+-----+
|      0 |      0 |      0 |      0 |
|      0 |      0 |      0 |      0 |
|      1 |      1 |      1 |      1 |
|      1 |      1 |      1 |      1 |
|      2 |      2 |      2 |      2 |
|      2 |      2 |      2 |      2 |
+-----+-----+-----+-----+
6 rows in set (3 disk I/O, 224.56 ms)

tinysql> DELETE FROM example WHERE attr0 = 1
Query OK (4 disk I/O, 299.08 ms)

tinysql> SELECT * FROM example
+-----+-----+-----+-----+
| attr0 | attr1 | attr2 | attr3 |
+-----+-----+-----+-----+
|      0 |      0 |      0 |      0 |
|      0 |      0 |      0 |      0 |
|      2 |      2 |      2 |      2 |
|      2 |      2 |      2 |      2 |
+-----+-----+-----+-----+
4 rows in set (2 disk I/O, 149.34 ms)
```

5.3 Selection

TinySQL will decompose the conditions joined by AND operation, and try to push down them to the bottom of the logic query plan. If a selection operation only involve a give table, it will carry out immediately after the scan of the table.

Example: Push down selection

```
tinysql> SELECT * FROM example, example1 WHERE example.attr0 = 0 AND
        example1.attr0 < 1

original logic query plan:
  '- WHERE
    |- WHERE_OPTION
    |   '- AND
    |     |- COMP_OP: =
    |     |   |- COLUMN_NAME: example.attr0
    |     |   '- INTEGER: 0
```

```

|           \- COMP_OP: <
|           |- COLUMN_NAME: example1.attr0
|           \- INTEGER: 1
\-- CROSS_PRODUCT
    |- BASE_NODE: example
    \- BASE_NODE: example1

optimized logic query plan:
\-- WHERE
    |- WHERE_OPTION
    |   \- AND
    |       |- COMP_OP: =
    |       |   |- COLUMN_NAME: example.attr0
    |       |   \- INTEGER: 0
    |       \- COMP_OP: <
    |           |- COLUMN_NAME: example1.attr0
    |           \- INTEGER: 1
    \-- CROSS_PRODUCT
        |- WHERE
        |   |- COMP_OP: =
        |   |   |- COLUMN_NAME: example.attr0
        |   |   \- INTEGER: 0
        |   \- BASE_NODE: example
        \-- WHERE
            |- COMP_OP: <
            |   |- COLUMN_NAME: example1.attr0
            |   \- INTEGER: 1
            \- BASE_NODE: example1

```

5.4 Join

Cross product with equation will be transformed to natural join. The disk I/O of natural join varies. In some worst case, it is possible that the overall disk I/Os increases (ex: all the tuple has the same joined attributes). However, most of the time, the time complexity will reduce from $O(B(R) \times B(S))$ of original cross product to $O(B(R) + B(S))$ of natural join.

For the pure cross product with more than 3 relations, it will do the product on the relations with the smallest number of tuples first.

Example: Transform cross product to natural join

```
tinysql> SELECT * FROM example, example1 WHERE example.attr0 = example1.attr0
```

original logic query plan:

```

\-- WHERE
    |- WHERE_OPTION
    |   \- COMP_OP: =
    |       |- COLUMN_NAME: example.attr0
    |       \- COLUMN_NAME: example1.attr0
    \-- CROSS_PRODUCT
        |- BASE_NODE: example
        \- BASE_NODE: example1

```

optimized logic query plan:

```
`- WHERE
  |- WHERE_OPTION
  |   `- COMP_OP: =
  |     |- COLUMN_NAME: example.attr0
  |     |- COLUMN_NAME: example1.attr0
  |- NATURAL_JOIN: {example.attr0, example1.attr0}
    |- BASE_NODE: example
    |- BASE_NODE: example1
```